

✓ Loading the dataset using Kaggle API

```
!pip install kaggle
```

Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.5.16)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (from kaggle) (2023.7.22)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.31.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from kaggle) (4.66.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (from kaggle) (8.0.1)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.0.7)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from kaggle) (6.1.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.4)

```
# ! mkdir ~/.kaggle
```

mkdir: cannot create directory ‘/root/.kaggle’: File exists

```
# ! cp kaggle.json ~/.kaggle/
```

```
# ! chmod 600 ~/.kaggle/kaggle.json
```

chmod: cannot access '/root/kaggle/kaggle.json': No such file or directory

```
from google.colab import files
```

```
files.upload()
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

{'kaggle.json': b'{"username": "shauryadhoundiyal", "key": "79ac760dd8639c19ed9678f7a4553226"}'}

```
! chmod 600 ~/.kaggle/kaggle.json
```

```
! kaggle datasets list
```

| ref | title | size |
|--|--|-------|
| carlmcbrideellis/llm-7-prompt-training-dataset | LLM: 7 prompt training dataset | 41MB |
| thedrcat/daigt-proper-train-dataset | DAIGT Proper Train Dataset | 119MB |
| iamsouravbanerjee/customer-shopping-trends-dataset | Customer Shopping Trends Dataset | 146KB |
| joebeachcapital/30000-spotify-songs | 30000 Spotify Songs | 3MB |
| thedrcat/daigt-v2-train-dataset | DAIGT V2 Train Dataset | 29MB |
| prasad22/healthcare-dataset | Healthcare Dataset | 48KB |
| mauryansshivam/list-of-internet-products-of-top-tech-companies | List of Internet Products of Top Tech Companies | 9KB |
| dillonmyrick/high-school-student-performance-and-demographics | High School Student Performance & Demographics | 24KB |
| nelgiriyeewithana/world-educational-data | World Educational Data | 9KB |
| alejopaullier/daigt-external-dataset | DAIGT External Dataset | 3MB |
| joebeachcapital/coronavirus-covid-19-cases-daily-updates | Coronavirus (COVID-19) Cases (Daily Updates) | 14MB |
| jacksondivakarr/online-shopping-dataset | Online Shopping Dataset | 146KB |
| anshtanwar/top-200-trending-books-with-reviews | Top 100 Bestselling Book Reviews on Amazon | 422KB |
| stefancomanita/top-us-songs-from-1950-to-2019-w-lyrics | Top US Songs from 1950 to 2019, w. lyrics | 674KB |
| asaniczka/ufc-fighters-statistics | UFC Fighters' Statistics Dataset | 148KB |
| zeesolver/consumer-behavior-and-shopping-habits-dataset | Consumer Behavior and Shopping Habits Dataset: | 146KB |
| ddosad/auto-sales-data | Automobile Sales data | 79KB |
| jdaustralia/icc-cwc23-all-innings-cleaned | ICC Cricket World Cup CWC23 All innings | 28KB |
| vikramrn/icc-mens-odi-world-cup-wc-2023 | ICC mens cricket odi world cup wc 2023 - batting | 16KB |
| simonezappatini/body-fat-extended-dataset | Body Fat Extended Dataset | 12KB |

```
!kaggle datasets download -d snginh/teethdecay
```

teethdecay.zip: Skipping, found more recently modified local copy (use --force to force download)

```
!unzip /content/teethdecay.zip
```

inflating: teeth_dataset/train/caries/wc41_8.jpg
inflating: teeth_dataset/train/caries/wc41_9.jpg
inflating: teeth_dataset/train/caries/wc42.jpg
inflating: teeth_dataset/train/caries/wc42_0.jpg
inflating: teeth_dataset/train/caries/wc42_1.jpg
inflating: teeth_dataset/train/caries/wc42_10.jpg
inflating: teeth_dataset/train/caries/wc42_11.jpg
inflating: teeth_dataset/train/caries/wc42_12.jpg
inflating: teeth_dataset/train/caries/wc42_13.jpg
inflating: teeth_dataset/train/caries/wc42_14.jpg
inflating: teeth_dataset/train/caries/wc42_15.jpg
inflating: teeth_dataset/train/caries/wc42_16.jpg
inflating: teeth_dataset/train/caries/wc42_17.jpg
inflating: teeth_dataset/train/caries/wc42_18.jpg
inflating: teeth_dataset/train/caries/wc42_19.jpg
inflating: teeth_dataset/train/caries/wc42_2.jpg
inflating: teeth_dataset/train/caries/wc42_3.jpg
inflating: teeth_dataset/train/caries/wc42_4.jpg
inflating: teeth_dataset/train/caries/wc42_5.jpg
inflating: teeth_dataset/train/caries/wc42_6.jpg
inflating: teeth_dataset/train/caries/wc42_7.jpg
inflating: teeth_dataset/train/caries/wc42_8.jpg
inflating: teeth_dataset/train/caries/wc42_9.jpg
inflating: teeth_dataset/train/caries/wc43.jpg
inflating: teeth_dataset/train/caries/wc43_0.jpg
inflating: teeth_dataset/train/caries/wc43_1.jpg
inflating: teeth_dataset/train/caries/wc43_10.jpg
inflating: teeth_dataset/train/caries/wc43_11.jpg
inflating: teeth_dataset/train/caries/wc43_12.jpg
inflating: teeth_dataset/train/caries/wc43_13.jpg
inflating: teeth_dataset/train/caries/wc43_14.jpg
inflating: teeth_dataset/train/caries/wc43_15.jpg
inflating: teeth_dataset/train/caries/wc43_16.jpg
inflating: teeth_dataset/train/caries/wc43_17.jpg
inflating: teeth_dataset/train/caries/wc43_18.jpg
inflating: teeth_dataset/train/caries/wc43_19.jpg
inflating: teeth_dataset/train/caries/wc43_2.jpg
inflating: teeth_dataset/train/caries/wc43_3.jpg
inflating: teeth_dataset/train/caries/wc43_4.jpg
inflating: teeth_dataset/train/caries/wc43_5.jpg
inflating: teeth_dataset/train/caries/wc43_6.jpg
inflating: teeth_dataset/train/caries/wc43_7.jpg
inflating: teeth_dataset/train/caries/wc43_8.jpg
inflating: teeth_dataset/train/caries/wc43_9.jpg
inflating: teeth_dataset/train/caries/wc47.jpg
inflating: teeth_dataset/train/caries/wc47_0.jpg
inflating: teeth_dataset/train/caries/wc47_1.jpg
inflating: teeth_dataset/train/caries/wc47_10.jpg
inflating: teeth_dataset/train/caries/wc47_11.jpg
inflating: teeth_dataset/train/caries/wc47_12.jpg
inflating: teeth_dataset/train/caries/wc47_13.jpg

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tqdm import tqdm
import cv2

import warnings
warnings.filterwarnings('ignore')
import os
train_path = "/content/teeth_dataset/train/"
test_path = "/content/teeth_dataset/test"
```


```
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Activation, Dropout, Flatten, Dense, BatchNormalization
from keras.preprocessing.image import ImageDataGenerator
from keras.utils import img_to_array, load_img
from keras.utils import plot_model
from glob import glob
```


✖ Getting contents of training data

```
x_data = []
y_data = []


for category in glob(train_path+'/*'):
    for file in tqdm(glob(category+'/*')):
        img_array=cv2.imread(file)
        img_array = cv2.cvtColor(img_array, cv2.COLOR_BGR2RGB)
        x_data.append(img_array)
        y_data.append(category.split("/")[-1])

data=pd.DataFrame({'image': x_data,'label': y_data})
```


 100%|██████████| 945/945 [00:00<00:00, 1429.97it/s]

 100%|██████████| 315/315 [00:00<00:00, 1382.21it/s]


```
data.shape
```

 (1260, 2)

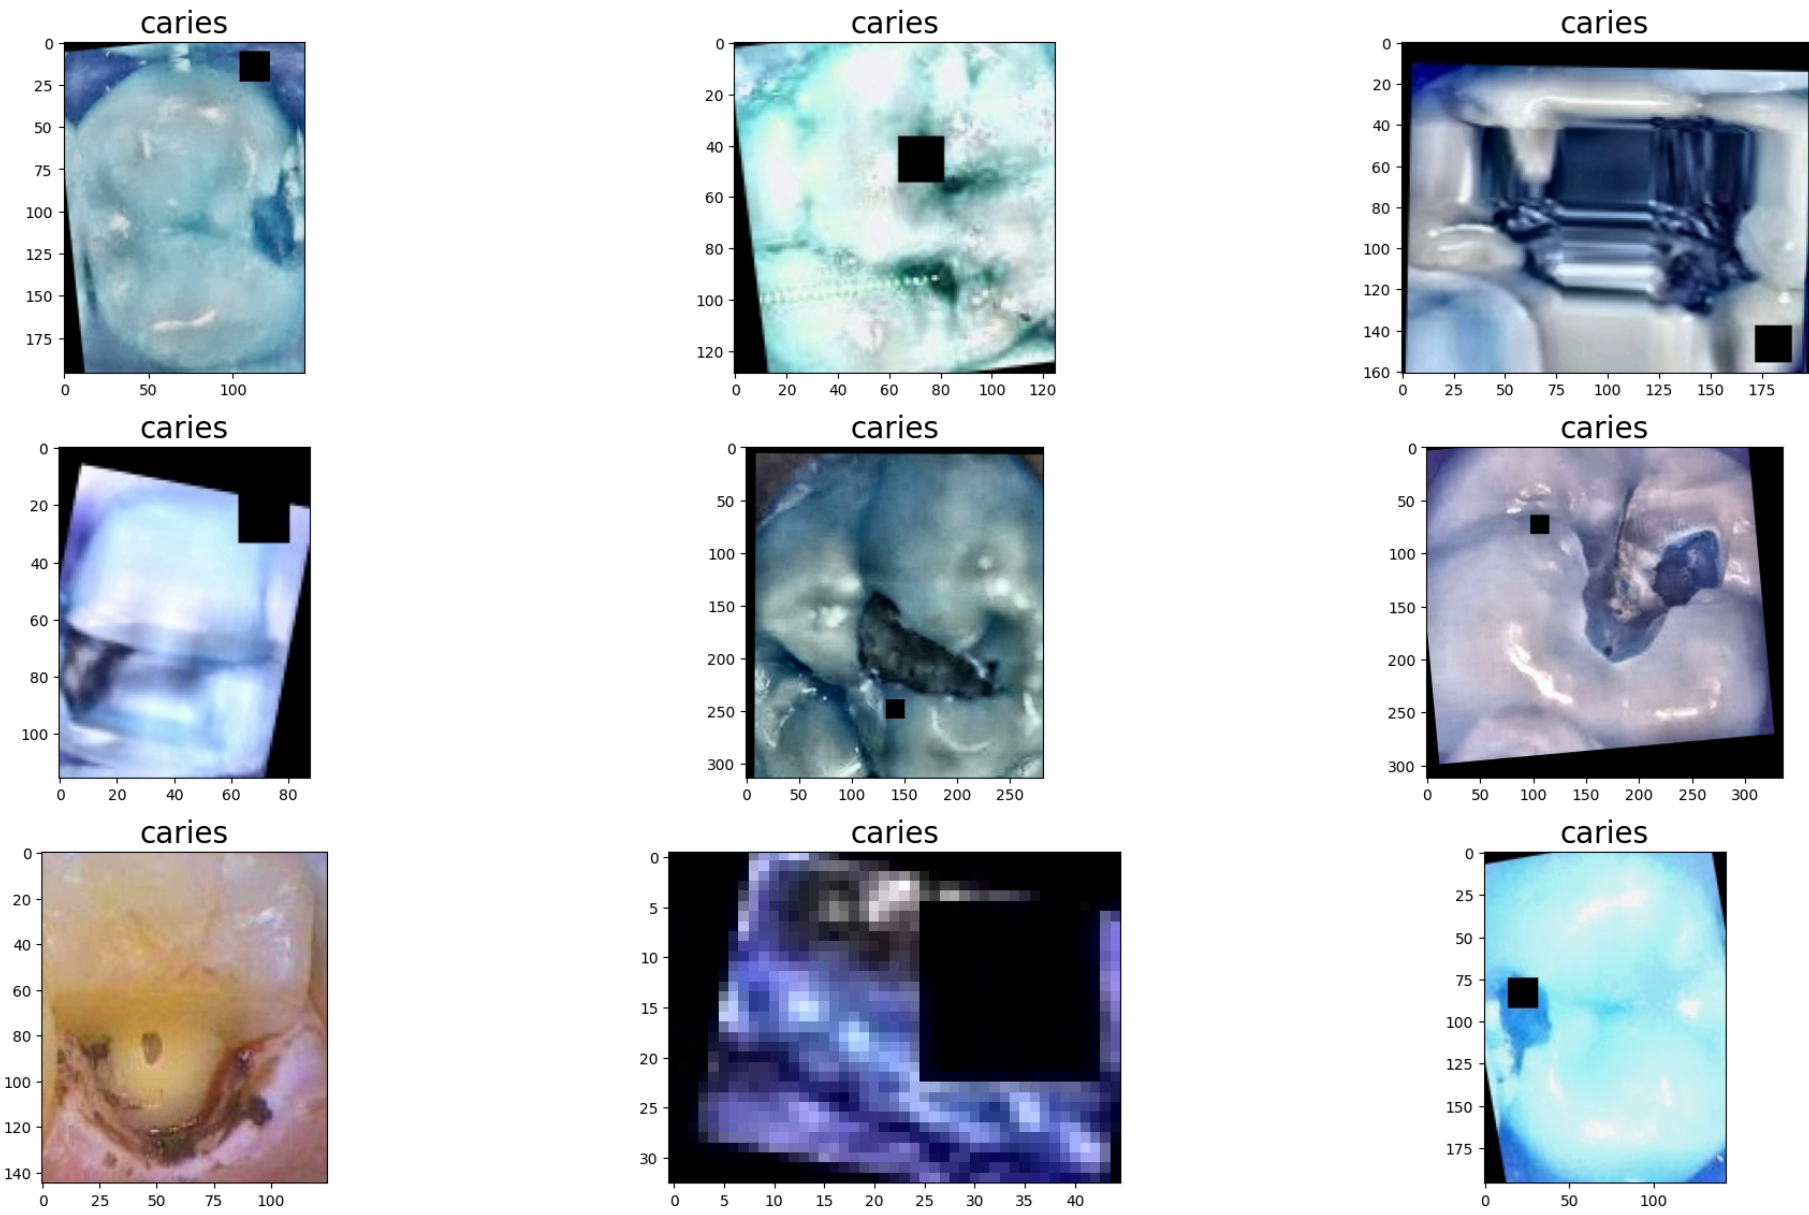
```
from collections import Counter
Counter(y_data)
```

 Counter({'caries': 945, 'no-caries': 315})

```
colors = ['green','blue']
plt.pie(data.label.value_counts(),startangle=90,explode=[0.05,0.05],autopct='%0.2f%%',labels=['No Caries', 'Caries'], color
plt.show()
```



```
plt.figure(figsize=(20,15))
for i in range(9):
    plt.subplot(4,3,(i%12)+1)
    index=np.random.randint(1000)
    plt.title('{0}'.format(data.label[index]),fontdict={'size':20})
    plt.imshow(data.image[index])
    plt.tight_layout()
```



```
className = glob(train_path + '/*' )
numberOfClass = len(className)
print("Number Of Class: ",numberOfClass)
```



Number Of Class: 2

Gettings contents of testing data

```
x_data = []
y_data = []

for category in glob(test_path+'/*'):
    for file in tqdm(glob(category+'/*')):
        img_array=cv2.imread(file)
        img_array = cv2.cvtColor(img_array, cv2.COLOR_BGR2RGB)
        x_data.append(img_array)
        y_data.append(category.split("/")[-1])

data=pd.DataFrame({'image': x_data,'label': y_data})
```



100%|██████████| 228/228 [00:00<00:00, 1530.77it/s]
100%|██████████| 84/84 [00:00<00:00, 4802.18it/s]

```
data.shape
```




(312, 2)

```
from collections import Counter
Counter(y_data)
```

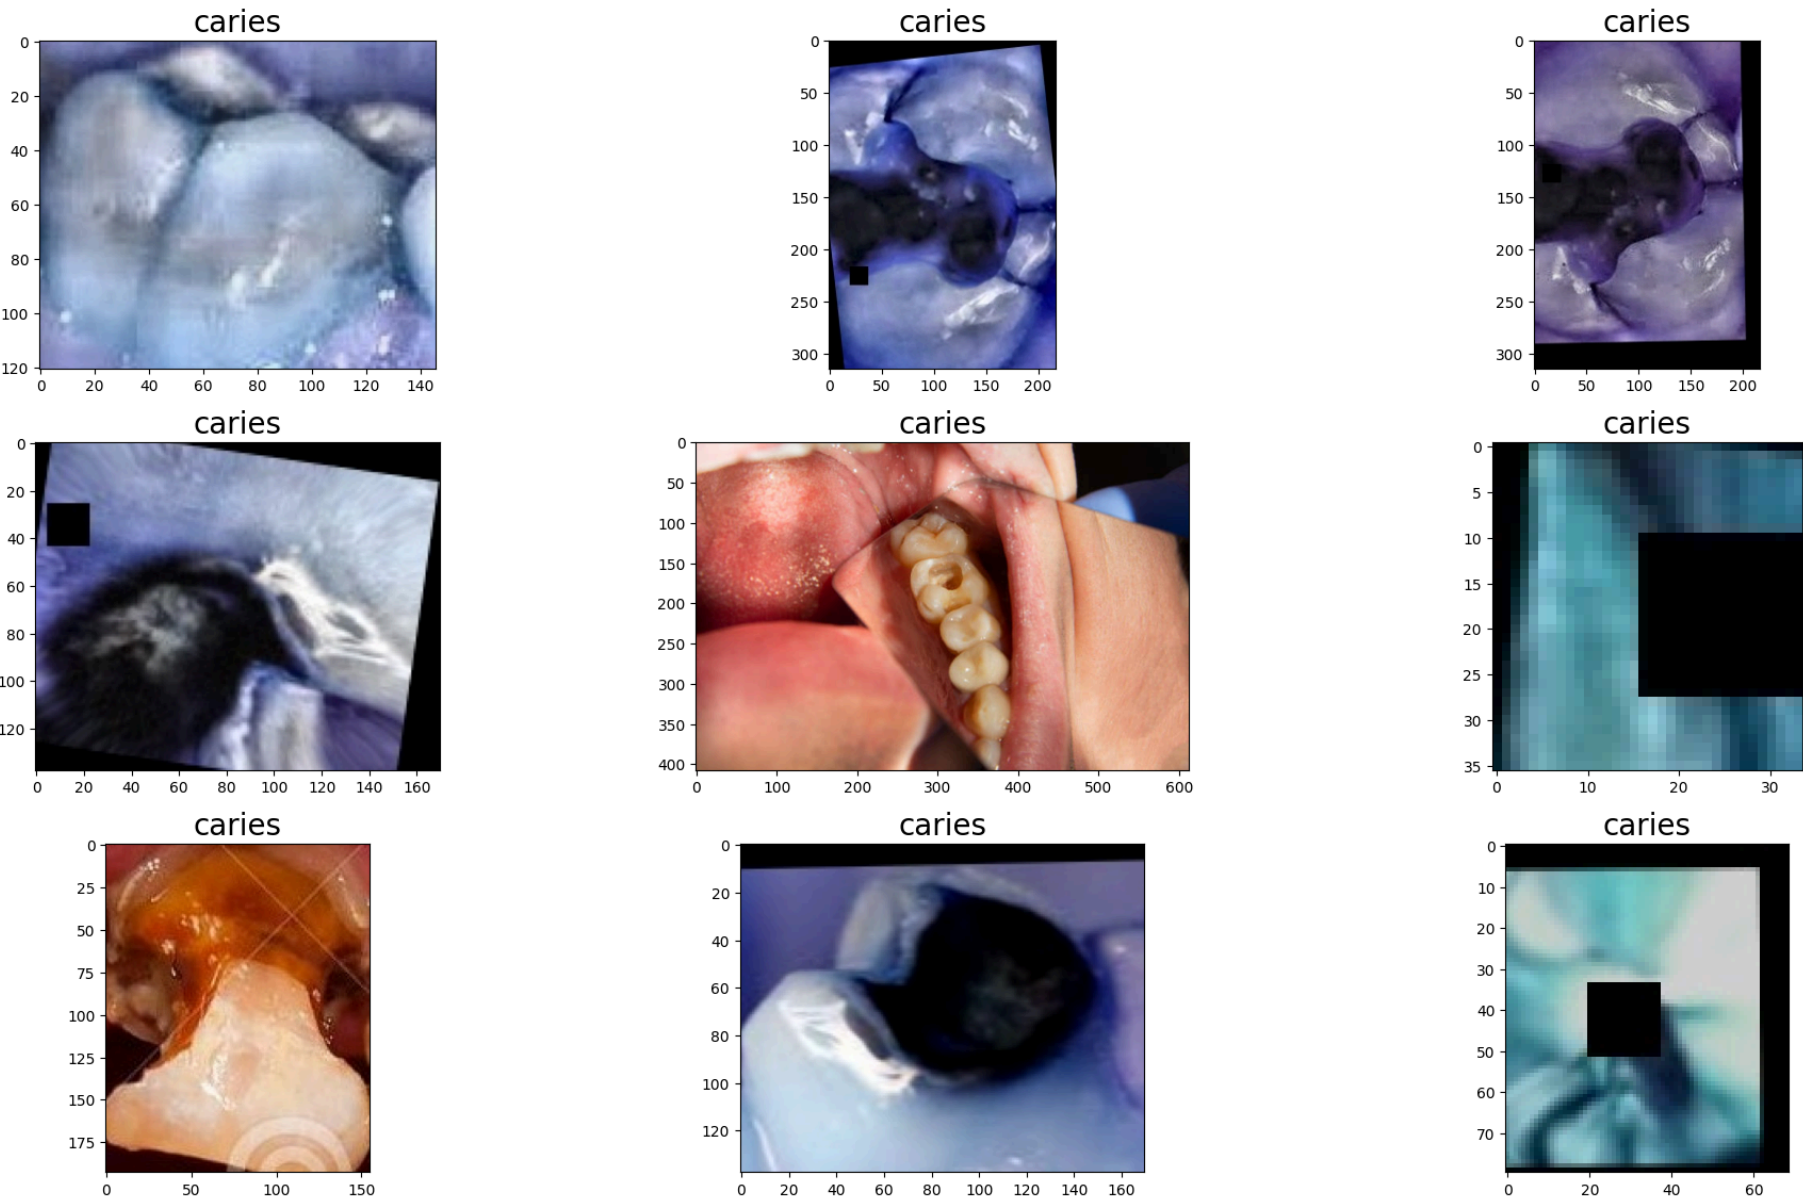


Counter({'caries': 228, 'no-caries': 84})

```
colors = ['green','blue']
plt.pie(data.label.value_counts(),startangle=90,explode=[0.05,0.05],autopct='%0.2f%%',labels=['No Caries', 'Caries'], color
plt.show()
```




```
plt.figure(figsize=(20,15))
for i in range(9):
    plt.subplot(4,3,(i%12)+1)
    index=np.random.randint(294)
    plt.title('{0}'.format(data.label[index]),fontdict={'size':20})
    plt.imshow(data.image[index])
    plt.tight_layout()
```



Model Building


```
train_generator = ImageDataGenerator(
    rescale = 1.0/255.,
    rotation_range=30,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    vertical_flip=True,
    fill_mode='nearest',
    brightness_range=(0.5, 1.5),
)
test_generator = ImageDataGenerator(rescale=1./255.)
```

```
train_datagen = train_generator.flow_from_directory(
    train_path,
    batch_size=10,
    target_size=(224, 224)
)
test_datagen = test_generator.flow_from_directory(
    test_path,
    batch_size=10,
    target_size=(224, 224)
)
```

 Found 1260 images belonging to 2 classes.
Found 312 images belonging to 2 classes.

Proposed CNN Model

```
model = Sequential()

model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))


model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))

model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(2, activation='softmax'))

model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
print(model.summary())
```

 Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------|----------------------|---------|
| ===== | | |
| conv2d (Conv2D) | (None, 222, 222, 32) | 896 |
| max_pooling2d (MaxPooling2D) | (None, 111, 111, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 109, 109, 64) | 18496 |
| max_pooling2d_1 (MaxPooling2D) | (None, 54, 54, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 52, 52, 128) | 73856 |
| max_pooling2d_2 (MaxPooling2D) | (None, 26, 26, 128) | 0 |
| conv2d_3 (Conv2D) | (None, 24, 24, 128) | 147584 |
| max_pooling2d_3 (MaxPooling2D) | (None, 12, 12, 128) | 0 |

7/20/24, 11:42 PM

DentalCavityClassification.ipynb - Colab

| | | |
|--------------------------------------|---------------|---------|
| flatten (Flatten) | (None, 18432) | 0 |
| dense (Dense) | (None, 512) | 9437696 |
| dropout (Dropout) | (None, 512) | 0 |
| dense_1 (Dense) | (None, 2) | 1026 |
| ===== | | |
| Total params: 9679554 (36.92 MB) | | |
| Trainable params: 9679554 (36.92 MB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |
| <hr/> | | |
| None | | |

```
hist_cnn = model.fit_generator(generator = train_datagen,epochs=20,validation_data = test_datagen)
```

Epoch 1/20
126/126 [=====] - 36s 182ms/step - loss: 0.5945 - accuracy: 0.7468 - val_loss: 0.5579 - val_ac
Epoch 2/20
126/126 [=====] - 22s 177ms/step - loss: 0.5404 - accuracy: 0.7516 - val_loss: 0.3908 - val_ac
Epoch 3/20
126/126 [=====] - 22s 170ms/step - loss: 0.5094 - accuracy: 0.7643 - val_loss: 0.4236 - val_ac
Epoch 4/20
126/126 [=====] - 22s 173ms/step - loss: 0.5021 - accuracy: 0.7698 - val_loss: 0.4083 - val_ac
Epoch 5/20
126/126 [=====] - 20s 162ms/step - loss: 0.4899 - accuracy: 0.7706 - val_loss: 0.3811 - val_ac
Epoch 6/20
126/126 [=====] - 22s 170ms/step - loss: 0.4978 - accuracy: 0.7706 - val_loss: 0.3968 - val_ac
Epoch 7/20
126/126 [=====] - 20s 162ms/step - loss: 0.4514 - accuracy: 0.7960 - val_loss: 0.3228 - val_ac
Epoch 8/20
126/126 [=====] - 22s 171ms/step - loss: 0.4467 - accuracy: 0.7897 - val_loss: 0.3037 - val_ac
Epoch 9/20
126/126 [=====] - 22s 172ms/step - loss: 0.4295 - accuracy: 0.7921 - val_loss: 0.3349 - val_ac
Epoch 10/20
126/126 [=====] - 20s 160ms/step - loss: 0.4247 - accuracy: 0.7960 - val_loss: 0.3092 - val_ac
Epoch 11/20
126/126 [=====] - 22s 173ms/step - loss: 0.3938 - accuracy: 0.8032 - val_loss: 0.4001 - val_ac
Epoch 12/20
126/126 [=====] - 21s 168ms/step - loss: 0.4048 - accuracy: 0.8238 - val_loss: 0.3729 - val_ac
Epoch 13/20
126/126 [=====] - 20s 159ms/step - loss: 0.4033 - accuracy: 0.8103 - val_loss: 0.3210 - val_ac
Epoch 14/20
126/126 [=====] - 21s 167ms/step - loss: 0.3743 - accuracy: 0.8230 - val_loss: 0.2927 - val_ac
Epoch 15/20
126/126 [=====] - 20s 158ms/step - loss: 0.4101 - accuracy: 0.8071 - val_loss: 0.3475 - val_ac
Epoch 16/20
126/126 [=====] - 22s 174ms/step - loss: 0.3721 - accuracy: 0.8270 - val_loss: 0.4127 - val_ac
Epoch 17/20
126/126 [=====] - 22s 177ms/step - loss: 0.3906 - accuracy: 0.8214 - val_loss: 0.3309 - val_ac
Epoch 18/20
126/126 [=====] - 20s 162ms/step - loss: 0.3764 - accuracy: 0.8190 - val_loss: 0.3634 - val_ac
Epoch 19/20
126/126 [=====] - 21s 169ms/step - loss: 0.3696 - accuracy: 0.8357 - val_loss: 0.3431 - val_ac
Epoch 20/20
126/126 [=====] - 20s 155ms/step - loss: 0.3797 - accuracy: 0.8214 - val_loss: 0.3381 - val_ac

Transfer Learning

ResNet50

```
from tensorflow.keras.layers import Dense, Flatten, Dropout, GlobalAveragePooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.applications.resnet50 import ResNet50


base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze pre-trained layers
for layer in base_model.layers:
    layer.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
predictions = Dense(2, activation='softmax')(x)

model_res = Model(inputs=base_model.input, outputs=predictions)

model_res.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
hist_res = model_res.fit(train_datagen, epochs=25, validation_data=test_datagen)
```

 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_order

94765736/94765736 [=====] - 3s 0us/step

| | |
|-------------|--|
| Epoch 1/25 | 126/126 [=====] - 32s 196ms/step - loss: 0.8482 - accuracy: 0.5897 - val_loss: 0.5780 - val_ac |
| Epoch 2/25 | 126/126 [=====] - 22s 175ms/step - loss: 0.7010 - accuracy: 0.6730 - val_loss: 0.5908 - val_ac |
| Epoch 3/25 | 126/126 [=====] - 23s 186ms/step - loss: 0.6163 - accuracy: 0.7151 - val_loss: 0.5290 - val_ac |
| Epoch 4/25 | 126/126 [=====] - 24s 188ms/step - loss: 0.6213 - accuracy: 0.7286 - val_loss: 0.6505 - val_ac |
| Epoch 5/25 | 126/126 [=====] - 23s 186ms/step - loss: 0.5865 - accuracy: 0.7198 - val_loss: 0.5348 - val_ac |
| Epoch 6/25 | 126/126 [=====] - 23s 183ms/step - loss: 0.5570 - accuracy: 0.7167 - val_loss: 0.4607 - val_ac |
| Epoch 7/25 | 126/126 [=====] - 24s 190ms/step - loss: 0.5559 - accuracy: 0.7238 - val_loss: 0.4736 - val_ac |
| Epoch 8/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.5494 - accuracy: 0.7294 - val_loss: 0.4942 - val_ac |
| Epoch 9/25 | 126/126 [=====] - 23s 179ms/step - loss: 0.5342 - accuracy: 0.7373 - val_loss: 0.8879 - val_ac |
| Epoch 10/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.5159 - accuracy: 0.7476 - val_loss: 0.5141 - val_ac |
| Epoch 11/25 | 126/126 [=====] - 24s 187ms/step - loss: 0.5026 - accuracy: 0.7540 - val_loss: 0.9520 - val_ac |
| Epoch 12/25 | 126/126 [=====] - 24s 191ms/step - loss: 0.5235 - accuracy: 0.7500 - val_loss: 0.4913 - val_ac |
| Epoch 13/25 | 126/126 [=====] - 23s 184ms/step - loss: 0.5110 - accuracy: 0.7357 - val_loss: 0.8938 - val_ac |
| Epoch 14/25 | 126/126 [=====] - 22s 175ms/step - loss: 0.5093 - accuracy: 0.7500 - val_loss: 0.6056 - val_ac |
| Epoch 15/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.5086 - accuracy: 0.7460 - val_loss: 0.5095 - val_ac |
| Epoch 16/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.5168 - accuracy: 0.7484 - val_loss: 0.5191 - val_ac |
| Epoch 17/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.4984 - accuracy: 0.7484 - val_loss: 0.4967 - val_ac |
| Epoch 18/25 | 126/126 [=====] - 23s 181ms/step - loss: 0.4946 - accuracy: 0.7587 - val_loss: 0.5452 - val_ac |
| Epoch 19/25 | 126/126 [=====] - 24s 187ms/step - loss: 0.4948 - accuracy: 0.7429 - val_loss: 0.5021 - val_ac |
| Epoch 20/25 | 126/126 [=====] - 23s 185ms/step - loss: 0.4823 - accuracy: 0.7571 - val_loss: 1.1181 - val_ac |
| Epoch 21/25 | 126/126 [=====] - 24s 192ms/step - loss: 0.4992 - accuracy: 0.7444 - val_loss: 0.4686 - val_ac |
| Epoch 22/25 | 126/126 [=====] - 22s 176ms/step - loss: 0.4802 - accuracy: 0.7595 - val_loss: 0.6353 - val_ac |
| Epoch 23/25 | 126/126 [=====] - 23s 184ms/step - loss: 0.5014 - accuracy: 0.7437 - val_loss: 0.4121 - val_ac |
| Epoch 24/25 | 126/126 [=====] - 23s 184ms/step - loss: 0.4899 - accuracy: 0.7532 - val_loss: 0.4362 - val_ac |
| Epoch 25/25 | 126/126 [=====] - 23s 187ms/step - loss: 0.4718 - accuracy: 0.7619 - val_loss: 1.0615 - val_ac |

▼ InceptionV3


```
from tensorflow.keras.layers import Dense, Flatten, Dropout, GlobalAveragePooling2D, MaxPooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.applications.inception_v3 import InceptionV3


base_model = InceptionV3(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze pre-trained layers
for layer in base_model.layers:
    layer.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
predictions = Dense(2, activation='softmax')(x)

model_inc = Model(inputs=base_model.input, outputs=predictions)

model_inc.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
hist_inc = model_inc.fit(train_datagen, epochs=10, validation_data=test_datagen)
```

 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf87910968/87910968 [=====] - 3s 0us/step

Epoch 1/10

126/126 [=====] - 32s 192ms/step - loss: 0.7902 - accuracy: 0.6556 - val_loss: 0.4694 - val_ac

Epoch 2/10

126/126 [=====] - 23s 182ms/step - loss: 0.5961 - accuracy: 0.7484 - val_loss: 0.4000 - val_ac

Epoch 3/10

126/126 [=====] - 22s 174ms/step - loss: 0.5108 - accuracy: 0.7952 - val_loss: 0.3656 - val_ac

Epoch 4/10

126/126 [=====] - 22s 171ms/step - loss: 0.4440 - accuracy: 0.8190 - val_loss: 0.3710 - val_ac

Epoch 5/10

126/126 [=====] - 23s 186ms/step - loss: 0.4232 - accuracy: 0.8302 - val_loss: 0.3627 - val_ac

Epoch 6/10

126/126 [=====] - 23s 185ms/step - loss: 0.3946 - accuracy: 0.8341 - val_loss: 0.4063 - val_ac

Epoch 7/10

126/126 [=====] - 23s 181ms/step - loss: 0.3953 - accuracy: 0.8357 - val_loss: 0.3490 - val_ac

Epoch 8/10


126/126 [=====] - 22s 172ms/step - loss: 0.3575 - accuracy: 0.8563 - val_loss: 0.3665 - val_ac


Epoch 9/10

126/126 [=====] - 23s 182ms/step - loss: 0.3865 - accuracy: 0.8365 - val_loss: 0.3392 - val_ac

Epoch 10/10

126/126 [=====] - 23s 180ms/step - loss: 0.3763 - accuracy: 0.8389 - val_loss: 0.3205 - val_ac





▼ MobileNetV2

```
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.applications.mobilenet_v2 import MobileNetV2


base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze pre-trained layers
for layer in base_model.layers:
    layer.trainable = False


x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
predictions = Dense(2, activation='softmax')(x)

model_mob = Model(inputs=base_model.input, outputs=predictions)


model_mob.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
hist_mob = model_mob.fit(train_datagen, epochs=20, validation_data=test_datagen)
```

 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf9406464/9406464 [=====] - 1s 0us/step

| |
|--|
| Epoch 1/20 |
| 126/126 [=====] - 27s 169ms/step - loss: 0.7409 - accuracy: 0.6778 - val_loss: 0.5836 - val_ac |
| Epoch 2/20 |
| 126/126 [=====] - 21s 165ms/step - loss: 0.5323 - accuracy: 0.7889 - val_loss: 0.6449 - val_ac |
| Epoch 3/20 |
| 126/126 [=====] - 22s 174ms/step - loss: 0.4804 - accuracy: 0.8024 - val_loss: 0.5564 - val_ac |
| Epoch 4/20 |
| 126/126 [=====] - 21s 167ms/step - loss: 0.4293 - accuracy: 0.8206 - val_loss: 0.4567 - val_ac |
| Epoch 5/20 |
| 126/126 [=====] - 20s 158ms/step - loss: 0.4025 - accuracy: 0.8278 - val_loss: 0.5284 - val_ac |
| Epoch 6/20 |
| 126/126 [=====] - 21s 166ms/step - loss: 0.3741 - accuracy: 0.8413 - val_loss: 0.4920 - val_ac |
| Epoch 7/20 |
| 126/126 [=====] - 20s 160ms/step - loss: 0.3810 - accuracy: 0.8278 - val_loss: 0.4104 - val_ac |
| Epoch 8/20 |
| 126/126 [=====] - 21s 163ms/step - loss: 0.3838 - accuracy: 0.8294 - val_loss: 0.4593 - val_ac |
| Epoch 9/20 |
| 126/126 [=====] - 21s 164ms/step - loss: 0.3690 - accuracy: 0.8365 - val_loss: 0.4396 - val_ac |
| Epoch 10/20 |
| 126/126 [=====] - 19s 152ms/step - loss: 0.3577 - accuracy: 0.8389 - val_loss: 0.4291 - val_ac |
| Epoch 11/20 |
| 126/126 [=====] - 21s 166ms/step - loss: 0.3434 - accuracy: 0.8444 - val_loss: 0.3694 - val_ac |
| Epoch 12/20 |
| 126/126 [=====] - 20s 156ms/step - loss: 0.3310 - accuracy: 0.8579 - val_loss: 0.4025 - val_ac |
| Epoch 13/20 |
| 126/126 [=====] - 21s 165ms/step - loss: 0.3393 - accuracy: 0.8524 - val_loss: 0.4365 - val_ac |
| Epoch 14/20 |
| 126/126 [=====] - 20s 158ms/step - loss: 0.3334 - accuracy: 0.8579 - val_loss: 0.5060 - val_ac |
| Epoch 15/20 |
| 126/126 [=====] - 21s 165ms/step - loss: 0.3041 - accuracy: 0.8667 - val_loss: 0.4830 - val_ac |
| Epoch 16/20 |
| 126/126 [=====] - 21s 169ms/step - loss: 0.3186 - accuracy: 0.8651 - val_loss: 0.4749 - val_ac |
| Epoch 17/20 |
| 126/126 [=====] - 21s 168ms/step - loss: 0.3232 - accuracy: 0.8683 - val_loss: 0.4169 - val_ac |
| Epoch 18/20 |
| 126/126 [=====] - 20s 162ms/step - loss: 0.3048 - accuracy: 0.8786 - val_loss: 0.3896 - val_ac |
| Epoch 19/20 |
| 126/126 [=====] - 21s 170ms/step - loss: 0.3340 - accuracy: 0.8444 - val_loss: 0.4108 - val_ac |
| Epoch 20/20 |
| 126/126 [=====] - 22s 171ms/step - loss: 0.3179 - accuracy: 0.8643 - val_loss: 0.4076 - val_ac |



✓ VGG16



```
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.applications.vgg16 import VGG16


base_model = VGG16(weights='imagenet', include_top=False, input_shape=(48, 48, 3))

# Freeze pre-trained layers
for layer in base_model.layers:
    layer.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(512, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(256, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
predictions = Dense(2, activation='softmax')(x)

model_vgg = Model(inputs=base_model.input, outputs=predictions)

model_vgg.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
hist_vgg = model_vgg.fit(train_datagen, epochs=20, validation_data=test_datagen)
```

 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_58889256/58889256 [=====] - 2s 0us/step

| |
|--|
| Epoch 1/20 |
| 126/126 [=====] - 28s 188ms/step - loss: 0.7371 - accuracy: 0.6762 - val_loss: 0.5275 - val_ac |
| Epoch 2/20 |
| 126/126 [=====] - 23s 184ms/step - loss: 0.5607 - accuracy: 0.7571 - val_loss: 0.4779 - val_ac |
| Epoch 3/20 |
| 126/126 [=====] - 23s 185ms/step - loss: 0.5006 - accuracy: 0.7913 - val_loss: 0.5728 - val_ac |
| Epoch 4/20 |
| 126/126 [=====] - 23s 184ms/step - loss: 0.4444 - accuracy: 0.8119 - val_loss: 0.4403 - val_ac |
| Epoch 5/20 |
| 126/126 [=====] - 23s 185ms/step - loss: 0.4375 - accuracy: 0.8040 - val_loss: 0.3854 - val_ac |
| Epoch 6/20 |
| 126/126 [=====] - 23s 186ms/step - loss: 0.4093 - accuracy: 0.8230 - val_loss: 0.4081 - val_ac |
| Epoch 7/20 |
| 126/126 [=====] - 22s 175ms/step - loss: 0.4139 - accuracy: 0.8159 - val_loss: 0.4466 - val_ac |
| Epoch 8/20 |
| 126/126 [=====] - 23s 183ms/step - loss: 0.3751 - accuracy: 0.8381 - val_loss: 0.3822 - val_ac |
| Epoch 9/20 |
| 126/126 [=====] - 23s 182ms/step - loss: 0.4043 - accuracy: 0.8246 - val_loss: 0.2917 - val_ac |
| Epoch 10/20 |
| 126/126 [=====] - 23s 184ms/step - loss: 0.3527 - accuracy: 0.8437 - val_loss: 0.3810 - val_ac |
| Epoch 11/20 |
| 126/126 [=====] - 23s 184ms/step - loss: 0.3905 - accuracy: 0.8294 - val_loss: 0.3375 - val_ac |
| Epoch 12/20 |
| 126/126 [=====] - 25s 199ms/step - loss: 0.3585 - accuracy: 0.8405 - val_loss: 0.3634 - val_ac |
| Epoch 13/20 |
| 126/126 [=====] - 25s 199ms/step - loss: 0.3489 - accuracy: 0.8524 - val_loss: 0.3904 - val_ac |
| Epoch 14/20 |
| 126/126 [=====] - 25s 200ms/step - loss: 0.3585 - accuracy: 0.8373 - val_loss: 0.4102 - val_ac |
| Epoch 15/20 |
| 126/126 [=====] - 25s 200ms/step - loss: 0.3710 - accuracy: 0.8357 - val_loss: 0.3459 - val_ac |
| Epoch 16/20 |
| 126/126 [=====] - 26s 203ms/step - loss: 0.3669 - accuracy: 0.8429 - val_loss: 0.3558 - val_ac |
| Epoch 17/20 |
| 126/126 [=====] - 25s 197ms/step - loss: 0.3590 - accuracy: 0.8349 - val_loss: 0.3482 - val_ac |
| Epoch 18/20 |
| 126/126 [=====] - 25s 198ms/step - loss: 0.3342 - accuracy: 0.8524 - val_loss: 0.3615 - val_ac |
| Epoch 19/20 |
| 126/126 [=====] - 24s 193ms/step - loss: 0.3584 - accuracy: 0.8444 - val_loss: 0.3832 - val_ac |
| Epoch 20/20 |
| 126/126 [=====] - 23s 183ms/step - loss: 0.3318 - accuracy: 0.8571 - val_loss: 0.3754 - val_ac |

Model Evaluation

```
# Creating holders to store the model performance results
Deep_Learning_Model = []
Training_accuracy = []
Validation_accuracy = []
Training_loss = []
Validation_loss = []

#function to call for storing the results
def storeResults(model, a,b,c,d):
    Deep_Learning_Model.append(model)
    Training_accuracy.append(round(a, 3))
    Validation_accuracy.append(round(b, 3))
    Training_loss.append(round(c, 3))
    Validation_loss.append(round(d, 3))
```


✓ Proposed CNN Model

✓ Accuracy Curve

```
acc = hist_cnn.history['accuracy']
val_acc = hist_cnn.history['val_accuracy']

acc1 = hist_cnn.history['accuracy'][-1]
val_acc1 = hist_cnn.history['val_accuracy'][-1]

plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('Proposed CNN Model Accuracy')
plt.show()
```



✓ Loss Curve


```
loss = hist_cnn.history['loss']
val_loss = hist_cnn.history['val_loss']

loss1 = hist_cnn.history['loss'][-1]
val_loss1 = hist_cnn.history['val_loss'][-1]

plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='lower right')
plt.ylabel('Loss')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('Proposed CNN Model Loss')
plt.show()

#storing the results. The below mentioned order of parameter passing is important.

storeResults('Proposed CNN Model',acc1,val_acc1,
            loss1,val_loss1)
```




✓ ResNet50

✓ Accuracy Curve

```
acc = hist_res.history['accuracy']
val_acc = hist_res.history['val_accuracy']

acc1 = hist_res.history['accuracy'][-1]
val_acc1 = hist_res.history['val_accuracy'][-1]

plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('ResNet50 Accuracy')
plt.show()
```



▼ Loss Curve


```
loss = hist_res.history['loss']
val_loss = hist_res.history['val_loss']

loss1 = hist_res.history['loss'][-1]
val_loss1 = hist_res.history['val_loss'][-1]

plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='lower right')
plt.ylabel('Loss')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('ResNet50 Loss')
plt.show()

#storing the results. The below mentioned order of parameter passing is important.

storeResults('ResNet50',acc1,val_acc1,
            loss1,val_loss1)
```




▼ InceptionV3

▼ Accuracy Curve


```
acc = hist_inc.history['accuracy']
val_acc = hist_inc.history['val_accuracy']

acc1 = hist_inc.history['accuracy'][-1]
val_acc1 = hist_inc.history['val_accuracy'][-1]

plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('InceptionV3 Accuracy')
plt.show()
```



▼ Loss Curve

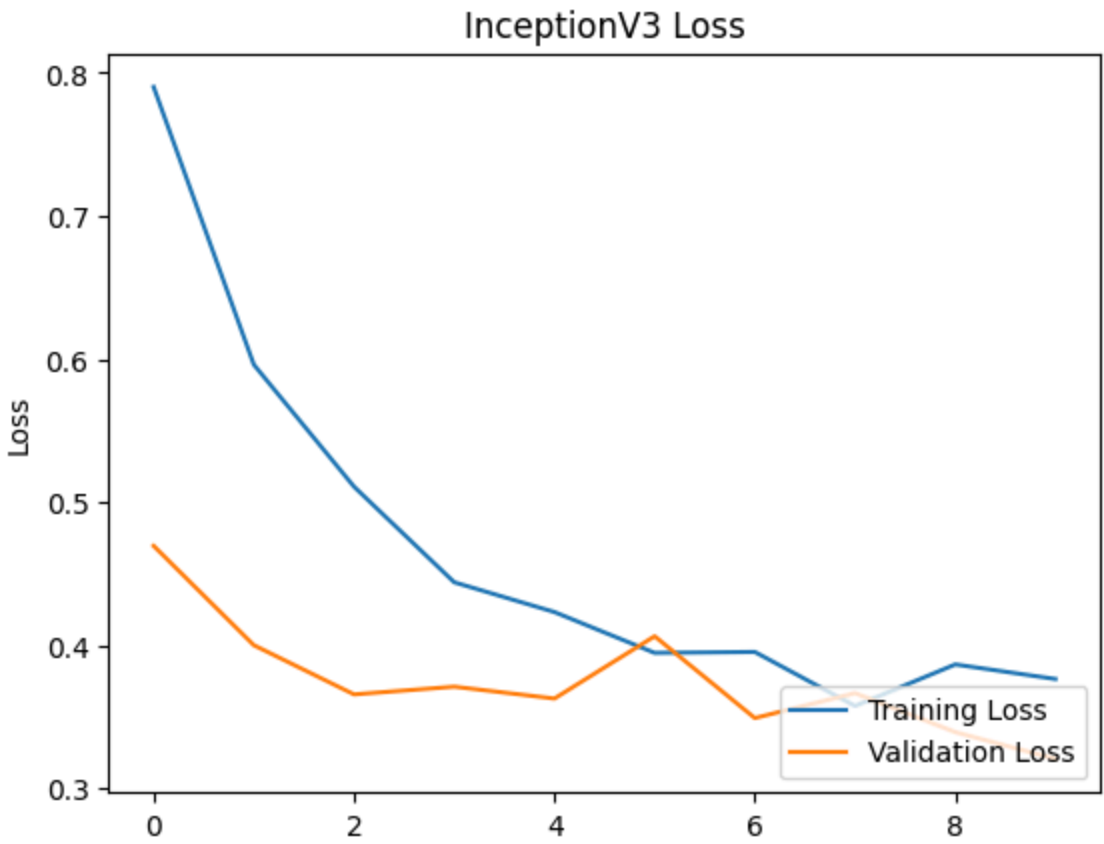
```
loss = hist_inc.history['loss']
val_loss = hist_inc.history['val_loss']

loss1 = hist_inc.history['loss'][-1]
val_loss1 = hist_inc.history['val_loss'][-1]

plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='lower right')
plt.ylabel('Loss')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('InceptionV3 Loss')
plt.show()

#storing the results. The below mentioned order of parameter passing is important.

storeResults('InceptionV3',acc1,val_acc1,
            loss1,val_loss1)
```




MobileNetV2

Accuracy Curve

```
acc = hist_mob.history['accuracy']
val_acc = hist_mob.history['val_accuracy']

acc1 = hist_mob.history['accuracy'][-1]
val_acc1 = hist_mob.history['val_accuracy'][-1]

plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('MobileNetV2 Accuracy')
plt.show()
```



Loss Curve


```
loss = hist_mob.history['loss']
val_loss = hist_mob.history['val_loss']

loss1 = hist_mob.history['loss'][-1]
val_loss1 = hist_mob.history['val_loss'][-1]

plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='lower right')
plt.ylabel('Loss')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
plt.title('MobileNetV2 Loss')
plt.show()

#storing the results. The below mentioned order of parameter passing is important.

storeResults('MobileNetV2',acc1,val_acc1,
            loss1,val_loss1)
```



✓ VGG16



✓ Accuracy Curve

```
acc = hist_vgg.history['accuracy']
val_acc = hist_vgg.history['val_accuracy']

acc1 = hist_vgg.history['accuracy'][-1]
val_acc1 = hist_vgg.history['val_accuracy'][-1]

plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),max(plt.ylim())])
```